

Invariance and Generalization in Sharpness-Aware Minimization (SAM)-Variants

Stanford CS229 Project

Manav Dahra

Department of Computer Science
Stanford University
mav3rik@stanford.edu

Abstract

In this study we train ResNet-18 & ViT-B/32 architectures under reparametrized settings using different variants of SAM optimizers. We compare their loss landscape curvature, generalisation gap and rate of convergence to find that overall, Monge-SAM performs the best across these metrics.

1 Introduction

1.1 Motivation

Recently, Sharpness-Aware Minimization (SAM) method and its variants have been proposed which use local perturbations at each optimization step to steer networks into flatter loss regions. Across these variants, the exact methodology of computing perturbations varies from relying on Euclidean metric for perturbation radius to using scale invariance and geometry-aware metrics. And even though these methods have shown theoretical and empirical improvements on generalization, they require a systematic empirical comparison under controlled settings. Additionally, each optimisation step in these methods require multiple forward and backward passes, which can be computationally expensive. It raises questions about the practical trade-offs between improved generalization and increased training time, especially for larger models and datasets.

This motivates our study to empirically validate some of these concepts by running extensive experiments on ResNet-18 and ViT-B/32 architectures using CIFAR-10 dataset as the input and predicted labels as the output. We focus on comparing how modern optimization paradigms, specifically SAM, Adaptive SAM (ASAM) and Monge SAM (M-SAM) actually withstand explicit parameter distortions against SGD optimizer.

2 Related Work

2.1 Background

The relationship between the geometry of the loss landscape and a model's generalization capacity has been a topic of interest in deep learning optimization research, fundamentally dating back to the concept of flat minima introduced by Hochreiter and Schmidhuber (1997) and expanded by Keskar et al. (2016). However, Dinh et al. (2017) demonstrated a critical flaw in traditional Euclidean metrics, proving that standard measures of flatness are highly sensitive to function-preserving reparametrizations. This was shown by doing node-rescaling in ReLU networks, meaning a minimum can be rendered arbitrarily sharp without altering its generalization behavior.

This insight revealed that the commonly held belief that flat minima inherently generalize better than sharp minima is not universally valid, as the notion of flatness can be manipulated through reparametrization. This has significant implications for optimization algorithms that rely on Euclidean

metrics to find flat minima, as they may be misled by the geometry of the parameter space rather than the true generalization properties of the model.

To address this limitation and actively seek flatter regions, researchers have developed optimization algorithms over the years which can be broadly categorized into following paradigms:

Traditional Euclidean Sharpness-Aware Methods (SAM):

Foret et al. (2021) introduced Sharpness-Aware Minimization (SAM), which augments SGD with a two-step update: parameters are first perturbed to the approximate local loss maximum within an ℓ_2 ball of radius ρ , then a standard gradient step is taken at the perturbed point. The key idea is to minimise the worst-case loss within a small neighbourhood, avoiding the prohibitively expensive Hessian. Since the perturbation ball is Euclidean it is *not* invariant to function-preserving rescalings of the weights, inheriting the vulnerability identified by Dinh et al. (2017).

Scale-Invariant Sharpness-Aware Methods (ASAM):

To address the scale-sensitivity of SAM, Kwon et al. (2021) replaces the ℓ_2 ball with a parameter-scaled adaptive ball, making the perturbation invariant to element-wise weight rescaling. This is computationally lightweight, but the diagonal scaling operator $T_w = \text{diag}(|w| + \eta)$ assumes all parameters are statistically independent, which may not hold in practice.

Geometry-Aware Sharpness-Aware Methods (Monge SAM):

The current SOTA paradigm treats the parameter space as a curved Riemannian manifold. Jacobsen and Arvanitidis (2025) equip it with the *Monge metric* $G(\theta) = I_K + \nabla \ell(\theta) \nabla \ell(\theta)^\top$, induced by embedding the loss surface as a graph in $\mathbb{R}^{K+1}$. This captures full inter-parameter correlations via the loss gradient, without requiring a probabilistic model formulation. See Table 1 for the resulting perturbation.

However, existing empirical validations of M-SAM are highly limited, restricted entirely to a narrow fine-tuning task on a pre-trained ResNet-18 and a multi-modal alignment task.

The Empirical Bridge: While these subsequent iterations offer robust theoretical guarantees for reparametrization invariance, existing empirical validations remain fragmented, lacking a systematically controlled reparametrization benchmark across diverse, modern neural architectures like Convolutional Networks and Vision Transformers. Our study unifies all three optimization paradigms into a highly controlled, multi-factorial grid—sweeping across architectures (ResNet-18 vs. ViT), initialization scales (1.0 vs. 5.0), perturbation radii ($\rho \in \{0.05, 0.1, 0.5\}$), and multiple random seeds (40, 41, 42). This comprehensive empirical framework allows us to directly compare the optimizers’ performance under explicit reparametrization distortions, providing a clearer picture of their practical robustness and generalization capabilities.

3 Dataset and Features

Dataset. We use CIFAR-10 (Krizhevsky, 2009), loaded via the HuggingFace Datasets library (Lhoest et al., 2021). It consists of 60,000 colour images across 10 balanced classes (airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck). We use the standard split: **50,000 training** and **10,000 test** images. No separate validation set is held out; early-stopping decisions are based on test accuracy after each epoch.

Input/Output. The input to each model is a pre-processed $3 \times H \times W$ image tensor; the output is a predicted class label from the 10 CIFAR-10 categories.

Preprocessing and augmentation. Raw images are 32×32 pixels. For ResNet-18 we keep the native resolution; for ViT-B/32 we upsample to 224×224 (the resolution expected by the pre-trained patch embedding). All images are normalised channel-wise using the per-channel empirical mean and standard deviation computed over the training set. During training we apply standard data augmentation: random crop with padding 4 and random horizontal flip. No augmentation is applied at test time beyond resizing and normalisation. See Appendix 7.1 for more details.

Features. We feed raw pixel values directly into the networks; no hand-crafted feature extraction (e.g. HOG, SIFT) is used. ResNet-18 learns hierarchical convolutional features from scratch on

CIFAR-10, while ViT-B/32 is fine-tuned from ImageNet pre-trained weights and processes the image as a sequence of 32×32 non-overlapping patches projected into a 768-dimensional embedding space.

Reparametrisation setup. To evaluate optimizer robustness, we introduce explicit weight distortions at initialisation by scaling all parameters by a factor $\alpha \in \{1.0, 5.0\}$. A scale of $\alpha = 1.0$ is the standard baseline; $\alpha = 5.0$ represents an adversarially sharp initialisation that exposes the scale-sensitivity of Euclidean perturbation methods.



Figure 1: Example images from the CIFAR-10 dataset (one per class). Classes are, from left to right: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck.

4 Methods

4.1 Experimental setup

Our experiment is split into following stages:

- **Baseline.** Run a standard training loop that sweeps across a grid of hyper-parameters without any reparametrization (i.e. $\alpha = 1.0$) to establish a baseline performance. For complete details of hyper-parameters see Appendix 7.2. The training loop also generates trained checkpoints and key metrics like loss, accuracy and elapsed time for all epochs and seeds.
- **Reparameterisation.** Next, we introduce reparametrization by scaling initial weights with $\alpha = 5.0$ and repeat the training loop for same hyper-parameter grid. This allows us to directly compare how each optimizer’s performance changes under explicit weight distortions, revealing their robustness to sharp initializations.
- **Analysis.** Next, we analyse collected results to compare optimizers across multiple dimensions: convergence speed, final test accuracy, generalization gap (train vs. test accuracy), and loss landscape curvature using Hutchinson’s trace estimation. We also visualise training curves and loss landscapes to qualitatively assess optimizer behavior under reparametrization.

4.2 Reparameterisation setup ResNet vs ViT

For ResNet-18 each BasicBlock has the following computation chain:

$$\dots \xrightarrow{\text{conv1}} \xrightarrow{\text{BN1}} \xrightarrow{\text{ReLU}} \xrightarrow{\text{conv2}} \xrightarrow{\times 1/\alpha} \xrightarrow{\text{BN2}} \dots$$

we can reparametrize it by injecting a positive scaling factor α at Batch normalisation layer (BN1) and compensating with $1/\alpha$ at Convolution layer (conv2), yielding the following chain which preserves the function computed by the block. See Appendix 7.2 for more details.

$$\dots \xrightarrow{\text{conv1}} \xrightarrow{\text{BN1}} \xrightarrow{\times \alpha} \xrightarrow{\text{ReLU}} \xrightarrow{\text{conv2}} \xrightarrow{\times 1/\alpha} \xrightarrow{\text{BN2}} \dots$$

For ViT-B/32, we apply reparameterisation in the MLP layers of the transformer blocks. Each transformer block contains multi-head self-attention and feedforward layers, with GELU activations. Since GELU activation used in ViT-B/32 is not positively homogeneous, meaning it does not satisfy $f(\alpha x) = \alpha f(x)$ for all $\alpha > 0$. This means that simply scaling the weights by α does not correspond to a function-preserving reparametrization, as the output of the GELU activation will not scale linearly with the input. To get around this issue, we patch the GELU activation function by its first-order Taylor expansion around zero, which is positively homogeneous. Specifically, we replace the GELU function with its linear approximation $f(x) \approx 0.5x$ for all inputs during training. This allows us to apply the same scaling factor α to the weights while maintaining a function-preserving reparametrization, enabling a fair comparison of optimizer performance under explicit weight distortions. To see how we arrived at this approximation, see Appendix 7.2.

4.3 Training loop and optimisers implementation

Since all the optimisers are variants of SGD, we implement a unified training loop that can accommodate the specific update rules for each optimizer.

We use out of the box implementation of SGD provided by PyTorch. While for SAM and ASAM, we build a wrapper around the SGD optimizer to implement the two-step update process as outlined in the original papers.

It is important to note that only the method of computing the perturbation $\hat{\epsilon}$ differs across the optimizers, as summarised in Table 1. Thus, we use a common codebase for the training loop, with conditional logic to compute $\hat{\epsilon}$ according to the specific optimizer being used. This ensures that all other aspects of the training process (e.g. learning rate, batch size, data augmentation) are held constant across optimizers, allowing for a fair comparison of their performance under reparametrization.

Table 1: Comparison of perturbation computation across SAM, ASAM, and M-SAM optimizers. Each optimizer defines a different constraint ball for the perturbation and a corresponding formula for computing $\hat{\epsilon}$ based on the loss gradient.

Optimizer	Constraint ball	Perturbation $\hat{\epsilon}$
SAM	$\ \hat{\epsilon}\ _2 \leq \rho$	$\frac{\rho}{\ \nabla_{\theta}\mathcal{L}(\theta)\ _2} \nabla_{\theta}\mathcal{L}(\theta)$
ASAM	$\ T_w^{-1}\hat{\epsilon}\ _2 \leq \rho$	$\frac{\rho T_w^2}{\ T_w \nabla_{\theta}\mathcal{L}(\theta)\ _2} \nabla_{\theta}\mathcal{L}(\theta)$
M-SAM	$\hat{\epsilon}^{\top} G(\theta) \hat{\epsilon} \leq \rho^2$	$\frac{\rho}{\sqrt{1 + \ \nabla_{\theta}\mathcal{L}(\theta)\ _2^2}} \nabla_{\theta}\mathcal{L}(\theta)$

4.4 Analysis details

To compare optimizers, we evaluate the following metrics for both baseline and reparametrized settings:

- **Training loss curves.** We plot training loss against epochs for each optimizer and α setting to visually assess convergence behavior and stability under reparametrization.
- **Convergence speed.** The number of epochs, GFLOPs and wall-clock time required to reach a certain accuracy threshold of 90, 95, 99% of the final test accuracy achieved by the baseline (SGD with $\alpha = 1.0$). This allows us to compare how quickly each optimizer converges under both standard and reparametrized settings.
- **Generalization gap.** The difference between training accuracy and test accuracy, indicating how well the model generalizes to unseen data.
- **Loss landscape curvature.** We use Hutchinson’s trace estimator to approximate the trace of the Hessian of the loss function at the final trained parameters, which serves as a measure of local curvature. A lower trace indicates a flatter minimum.

5 Results & Discussion

We present our results side by side between baseline and reparametrized settings to directly compare the impact of explicit weight distortions on each optimizer’s performance. We first discuss convergence speed, then generalization gap, and finally loss landscape curvature.

5.1 Training loss curves

During the training we capture epoch loss data, we use this data to plot loss curves and deduce some qualitative observations on convergence behavior and stability for each of the scenarios. For ResNet-18, we observe training loss curves trending downwards and converging to low values for both baseline and reparametrized settings. See Figure 6 for ResNet-18.

However, for ViT-B/32, the epoch loss curves exhibit high generalisation gap and slower convergence for SGD and SAM for the case (seed = 40 & $\rho = 0.1$). See Figure 2. This aligns with our expectations, for the reasons discussed above. From the figure, it is evident that A-SAM and M-SAM are more stable and faster in convergence compared to other optimizers.




Figure 2: Epoch-wise training loss curves for ViT-B/32 trained with SGD, SAM, ASAM, and M-SAM. The plot shows that M-SAM converges fastest and achieves a lower training loss compared to the other optimizers, indicating its effectiveness in finding flatter minima.

5.2 Convergence rate




Figure 3: shows 3 bar plots for *baseline* ResNet-18 trained with SGD(gray), SAM(blue), ASAM(orange), and M-SAM(green). Left most plot shows num of Epochs required to reach accuracy thresholds for 90%, 95%, 99% of the final training accuracy. Middle plot shows GFLOPs to reach the same accuracy thresholds. Rightmost plot shows wall-clock time to reach the same accuracy thresholds.




Figure 4: shows 3 bar plots for *baseline* ViT trained with SGD(gray), SAM(blue), ASAM(orange), and M-SAM(green). Left most plot shows num of Epochs required to reach accuracy thresholds for 90%, 95%, 99% of the final training accuracy. Middle plot shows GFLOPs to reach the same accuracy thresholds. Rightmost plot shows wall-clock time to reach the same accuracy thresholds.

Key observations from Figure 3 and Figure 4:

- Overall, M-SAM requires least number of epochs to reach the accuracy thresholds, indicating that its geometry-aware perturbation allows it to navigate the loss landscape more effectively than the other optimizers.
- SGD requires the least amount of GFLOPs, which is obvious since it performs the simplest updates. But, interestingly, M-SAM requires the second least amount of GFLOPs which can be attributed to the fact that it also requires least number of epochs to converge, thus offsetting the additional computational cost of its more complex perturbation.
- Wall clock time is more nuanced, with M-SAM being the fastest for ResNet-18 but not for ViT. This may be due to the fact that the computational overhead of M-SAM's perturbation is more significant for larger models like ViT, where the cost of computing the Monge metric and its inverse can become a bottleneck. However, even in ViT, M-SAM still outperforms ASAM and SGD in terms of wall-clock time, suggesting that its faster convergence still provides a net benefit despite the increased per-step cost.
- For simpler architecture (ResNet-18), SAM and M-SAM show similar convergence rates, both outperforming SGD and ASAM. Even though SGD takes least number of GFLOPs, it still requires most amount of epochs to converge to the final accuracy.
- For the more complex architecture (ViT-B/32), M-SAM shows a significantly faster convergence (at 99% accuracy) compared to all other optimizers, including SAM. This suggests that the geometry-aware perturbation in M-SAM may be particularly beneficial for navigating the more complex loss landscape of ViTs.
- Out of the two scale-invariant methods, ASAM shows a slower convergence compared to M-SAM, which may be due to its diagonal scaling operator not capturing inter-parameter correlations effectively. This is especially evident in the ViT architecture, where ASAM's performance is notably worse than M-SAM.

5.3 Generalization Gap




Figure 5: shows a bar plot comparing the average generalization gap (test loss - train loss) across all seeds for ResNet-18 trained with SGD, SAM, ASAM, and M-SAM under baseline ($\alpha = 1.0$) settings.

Key observations from Figure 5:

- SAM and M-SAM achieve the smallest generalization gap, indicating that they generalize best to unseen data. This suggests that M-SAM is effective at finding flatter minima that generalize well. It is interesting to see SAM performing comparably to M-SAM, which may be due to the fact that the ResNet-18 architecture is more stable and has smoother loss landscapes compared to larger models, making the Euclidean perturbation in SAM more effective.
- ASAM has a larger generalization gap than both SAM and M-SAM, indicating that its scale-invariant perturbation may not be as effective in improving generalization. This can be attributed to multiple factors, 1) the diagonal scaling operator in ASAM may not capture inter-parameter correlations, leading to suboptimal perturbations, and 2) the perturbation radius ρ values chosen for this experiment may not be in optimal range for ASAM. Further tuning of the hyper-parameters for ASAM could potentially improve its performance, but it does highlight to the fact that ASAM is more sensitive to hyper-parameter choices, which can be a practical limitation.
- SGD has the largest generalization gap, which is consistent with the expectation that it may converge to sharper minima that do not generalize as well to unseen data.

6 Conclusion / Future Work

Summarize your report and reiterate key points. Which algorithms were the highest-performing? Why do you think that some algorithms worked better than others? For future work, if you had more time, more team members, or more computational resources, what would you explore?

Note: All sections before this point must fit on five (5) pages. No exceptions. Supplemental material is not allowed. Anything else you want to add to your report (e.g. acknowledgements, author bios, funding sources) is included in the 5 page limit. The exception is the section describing the contributions of each team member. You will be penalized -10 points per page exceeding this limit. The max report score is 100.

7 Appendices

7.1 Preprocessing Details

Empirical mean and standard deviation for each channel are computed over the training set as follows:

$$\mu_c = \frac{1}{N} \sum_{i=1}^N x_{i,c} \quad \text{and} \quad \sigma_c = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_{i,c} - \mu_c)^2}$$

where N is the total number of training images, $x_{i,c}$ is the pixel value of the i -th image in channel c . The computed means and standard deviations are then used to normalise each image during preprocessing:

$$\hat{x}_{i,c} = \frac{x_{i,c} - \mu_c}{\sigma_c}$$

Values for the computed means and standard deviations for CIFAR-10 are:

- Red channel: $\mu_R = 0.4914$, $\sigma_R = 0.2470$
- Green channel: $\mu_G = 0.4822$, $\sigma_G = 0.2435$
- Blue channel: $\mu_B = 0.4465$, $\sigma_B = 0.2616$

7.2 Hyper-parameters

Table 2: Hyper-parameter grid used across all experiments. Baseline runs use $\alpha = 1.0$; reparametrization runs use $\alpha = 5.0$. Every combination of optimizer $\times \rho \times$ seed is evaluated independently, yielding $4 \times 3 \times 3 = 36$ runs per architecture per α setting.

Hyper-parameter	ResNet-18	ViT-B/32
Epochs	200	50
Batch size	256	256
Learning rate	0.1	0.01
Momentum	0.9	0.9
Weight decay	5×10^{-4}	5×10^{-4}
Input resolution	32×32	224×224
Pre-trained weights	No	Yes (ImageNet)
Optimizers	SGD, SAM, ASAM, M-SAM	
Perturbation radius ρ	{0.05, 0.1, 0.5}	
ASAM scaling η	0.01	
Random seeds	{40, 41, 42}	
Init. scale α (baseline)	1.0	
Init. scale α (reparam)	5.0	

7.3 Reparameterization

We need to be careful when reparametrising a layer before Batch normalisation (BN). BN absorbs any scaling factor applied to its input, meaning that if we simply scale the weights of a convolutional layer before BN by α , the output of BN will be unaffected, and the subsequent layers will not experience the intended reparameterization. To see this,

$$\mu' = \frac{1}{m} \sum_{i=1}^m \alpha \cdot x_i = \alpha \cdot \mu \quad \text{and} \quad \sigma' = \sqrt{\frac{1}{m} \sum_{i=1}^m (\alpha \cdot x_i - \alpha \cdot \mu)^2} = \alpha \cdot \sigma$$

$$\text{BN}(\alpha \cdot x) = \frac{\alpha \cdot x - \mu'}{\sigma'} = \frac{\alpha \cdot x - \alpha \cdot \mu}{\alpha \cdot \sigma} = \frac{x - \mu}{\sigma} = \text{BN}(x)$$

To get around this issue, we inject the scaling factor α at the output of BN1, which is then propagated through the ReLU activation. We compensate for this by applying a corresponding $1/\alpha$ scaling at conv2, ensuring that the overall function computed by the block remains unchanged.

Taylor approximation of GELU activation function around zero is given by:

$$\begin{aligned}\text{GELU}(x) &= 0.5x(1 + \tanh(\sqrt{2/\pi}(x + 0.044715x^3))) \\ &\approx 0.5x(1 + \tanh(\sqrt{2/\pi}x)) \\ &\approx 0.5x(1 + \sqrt{2/\pi}x) \\ &= 0.5x \quad \text{for small } x\end{aligned}$$

7.4 ResNet loss curves for baseline and reparametrized settings




Figure 6: Epoch-wise training loss curves for ResNet-18 trained with SGD, SAM, ASAM, and M-SAM. The plot shows that M-SAM converges faster and achieves a lower training loss compared to the other optimizers, indicating its effectiveness in finding flatter minima.

7.5 ViT loss curves for baseline and reparametrized settings




Figure 7: Epoch-wise training loss curves for ViT trained with SGD, SAM, ASAM, and M-SAM. The plot shows that M-SAM converges faster and achieves a lower training loss compared to the other optimizers, indicating its effectiveness in finding flatter minima.

8 Contributions

The contributions section is not included in the 5 page limit. This section should describe what each team member worked on and contributed to the project.

9 Note on citations

There are two citation commands:

Citation in parentheses `\citep{}` is for when you cite a paper at the end of a clause, and you could read the text out loud and not read the authors’ names. For example “We also run our experiments on a multilingual language model (?).”

Citation in text `\citet{}` is for in-text citations, when someone reading the text out loud would have to read the names of the authors for the sentence to make sense. For example, “We also run our experiments on the multilingual model described in ?.”

References

- Laurent Dinh, Razvan Pascanu, Samy Bengio, and Yoshua Bengio. 2017. Sharp minima can generalize for deep nets. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1019–1028. PMLR.
- Pierre Foret, Ariel Kleiner, Hossein Mobahi, and Behnam Neyshabur. 2021. Sharpness-aware minimization for efficiently improving generalization. In *International Conference on Learning Representations (ICLR)*.
- Sepp Hochreiter and Jürgen Schmidhuber. 1997. Flat minima. *Neural computation*, 9(1):1–42.

- Albert Kjøller Jacobsen and Georgios Arvanitidis. 2025. Monge SAM: Robust reparameterization-invariant sharpness-aware minimization based on loss geometry.
- Nitish Shirish Keskar, Dheevatsa Mudigere, Jorge Nocedal, Mikhail Smelyanskiy, and Ping Tak Peter Tang. 2016. On large-batch training for deep learning: Generalization gap and sharp minima. *arXiv preprint arXiv:1609.04836*.
- Alex Krizhevsky. 2009. Learning multiple layers of features from tiny images. Technical report, University of Toronto.
- Jungmin Kwon, Jeongseop Kim, Hyunseo Park, and In Kwon Choi. 2021. ASAM: Adaptive sharpness-aware minimization for scale-invariant learning of deep neural networks. In *International Conference on Machine Learning (ICML)*.
- Quentin Lhoest, Albert Villanova del Moral, Yacine Jernite, Abhishek Thakur, Patrick von Platen, Suraj Patil, Julien Chaumond, Mariama Drame, Julien Plu, Lewis Tunstall, Joe Davison, Mario Šaško, Gunjan Chhablani, Bhavitvya Malik, Simon Brandeis, Teven Le Scao, Victor Sanh, Canwen Xu, Nicolas Patry, Angelina McMillan-Major, Philipp Schmid, Sylvain Gugger, Clément Delangue, Théo Matussière, Lysandre Debut, Stas Bekman, Pierric Cistac, Thibault Goehringer, Victor Mustar, François Lagunas, Alexander Rush, and Thomas Wolf. 2021. Datasets: A community library for natural language processing. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 175–184, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.